

gender-and-age-prediction-cnn

January 31, 2024

0.1 Import Modules

```
[1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from tqdm.notebook import tqdm
warnings.filterwarnings('ignore')
%matplotlib inline

import tensorflow as tf
from keras.preprocessing.image import load_img
from keras.models import Sequential, Model
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D, Input
```

0.2 Load the Dataset

```
[2]: BASE_DIR = '../input/utkface-new/UTKFace/'

[3]: # labels - age, gender, ethnicity
image_paths = []
age_labels = []
gender_labels = []

for filename in tqdm(os.listdir(BASE_DIR)):
    image_path = os.path.join(BASE_DIR, filename)
    temp = filename.split('_')
    age = int(temp[0])
    gender = int(temp[1])
    image_paths.append(image_path)
    age_labels.append(age)
    gender_labels.append(gender)
```

```
0%|          | 0/23708 [00:00<?, ?it/s]
```

```
[4]: # convert to dataframe
df = pd.DataFrame()
df['image'], df['age'], df['gender'] = image_paths, age_labels, gender_labels
df.head()
```

```
[4]:
```

| | image | age | gender |
|---|---|-----|--------|
| 0 | ../input/utkface-new/UTKFace/26_0_2_2017010402... | 26 | 0 |
| 1 | ../input/utkface-new/UTKFace/22_1_1_2017011223... | 22 | 1 |
| 2 | ../input/utkface-new/UTKFace/21_1_3_2017010500... | 21 | 1 |
| 3 | ../input/utkface-new/UTKFace/28_0_0_2017011718... | 28 | 0 |
| 4 | ../input/utkface-new/UTKFace/17_1_4_2017010322... | 17 | 1 |

```
[5]: # map labels for gender
gender_dict = {0:'Male', 1:'Female'}
```


0.3 Exploratory Data Analysis

```
[6]: from PIL import Image
img = Image.open(df['image'][0])
plt.axis('off')
plt.imshow(img);
```




```
[7]: sns.distplot(df['age'])
```

```
[7]: <AxesSubplot:xlabel='age', ylabel='Density'>
```




```
[8]: sns.countplot(df['gender'])
```

```
[8]: <AxesSubplot:xlabel='gender', ylabel='count'>
```



```
[10]: # to display grid of images
plt.figure(figsize=(20, 20))
files = df.iloc[0:25]

for index, file, age, gender in files.itertuples():
    plt.subplot(5, 5, index+1)
    img = load_img(file)
    img = np.array(img)
    plt.imshow(img)
    plt.title(f"Age: {age} Gender: {gender_dict[gender]}")
    plt.axis('off')
```



0.4 Feature Extraction

```
[11]: def extract_features(images):
      features = []
      for image in tqdm(images):
          img = load_img(image, grayscale=True)
          img = img.resize((128, 128), Image.ANTIALIAS)
          img = np.array(img)
          features.append(img)

      features = np.array(features)
      # ignore this step if using RGB
      features = features.reshape(len(features), 128, 128, 1)
      return features
```

```
[13]: X = extract_features(df['image'])
```

```
0%|          | 0/23708 [00:00<?, ?it/s]
```

```
[15]: X.shape
```

```
[15]: (23708, 128, 128, 1)
```

```
[16]: # normalize the images
      X = X/255.0
```

```
[17]: y_gender = np.array(df['gender'])
      y_age = np.array(df['age'])
```

```
[18]: input_shape = (128, 128, 1)
```

0.5 Model Creation

```
[21]: inputs = Input((input_shape))
      # convolutional layers
      conv_1 = Conv2D(32, kernel_size=(3, 3), activation='relu')(inputs)
      maxp_1 = MaxPooling2D(pool_size=(2, 2))(conv_1)
      conv_2 = Conv2D(64, kernel_size=(3, 3), activation='relu')(maxp_1)
      maxp_2 = MaxPooling2D(pool_size=(2, 2))(conv_2)
      conv_3 = Conv2D(128, kernel_size=(3, 3), activation='relu')(maxp_2)
      maxp_3 = MaxPooling2D(pool_size=(2, 2))(conv_3)
      conv_4 = Conv2D(256, kernel_size=(3, 3), activation='relu')(maxp_3)
      maxp_4 = MaxPooling2D(pool_size=(2, 2))(conv_4)

      flatten = Flatten()(maxp_4)

      # fully connected layers
```

```
dense_1 = Dense(256, activation='relu') (flatten)
dense_2 = Dense(256, activation='relu') (flatten)

dropout_1 = Dropout(0.3) (dense_1)
dropout_2 = Dropout(0.3) (dense_2)

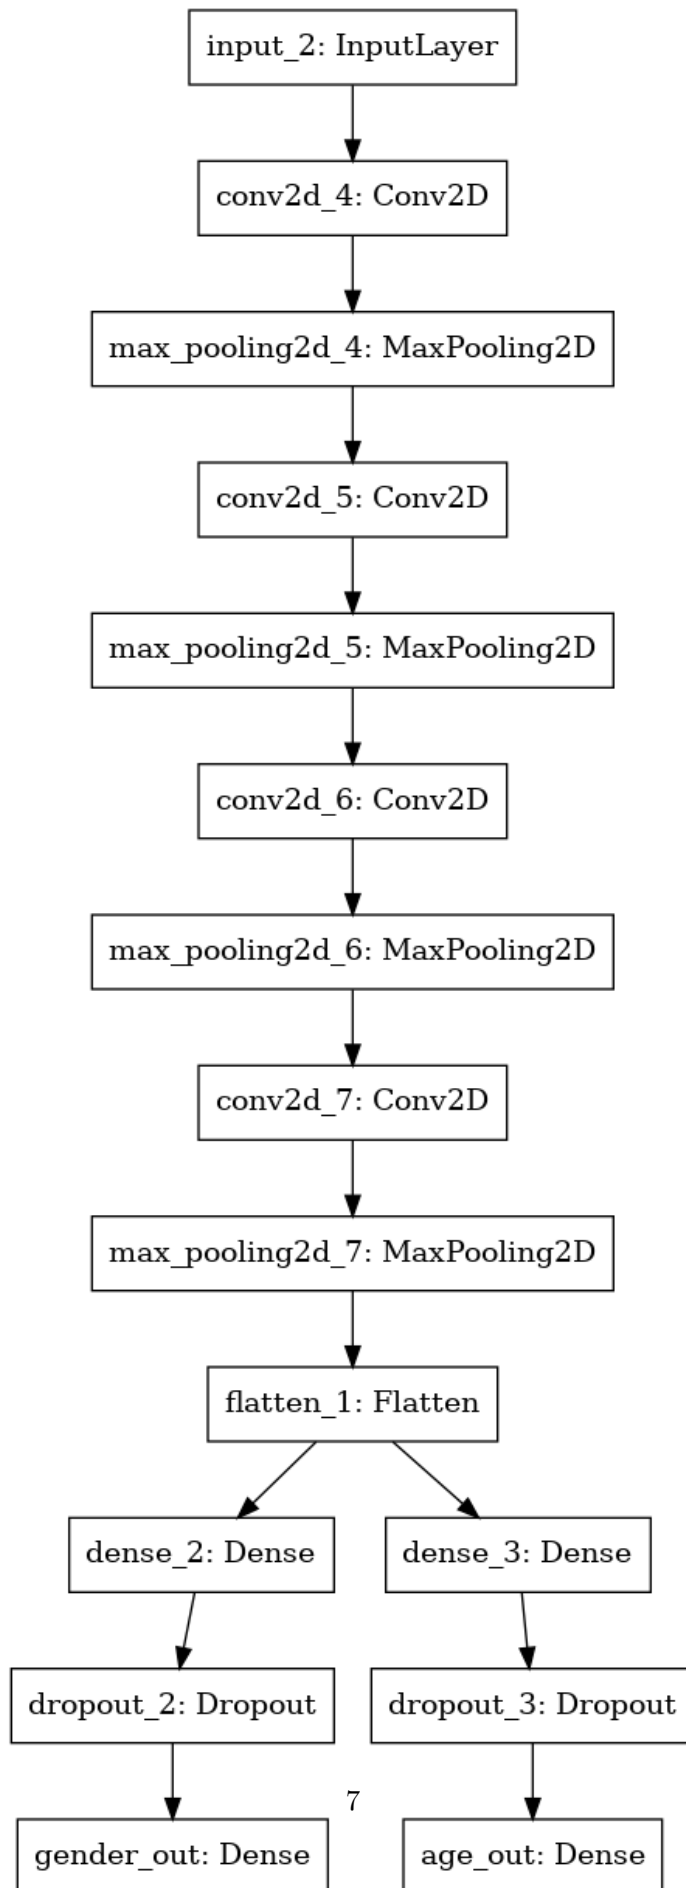
output_1 = Dense(1, activation='sigmoid', name='gender_out') (dropout_1)
output_2 = Dense(1, activation='relu', name='age_out') (dropout_2)

model = Model(inputs=[inputs], outputs=[output_1, output_2])

model.compile(loss=['binary_crossentropy', 'mae'], optimizer='adam',
               metrics=['accuracy'])
```

```
[25]: # plot the model
      from tensorflow.keras.utils import plot_model
      plot_model(model)
```

[25]:



```
[26]: # train model
history = model.fit(x=X, y=[y_gender, y_age], batch_size=32, epochs=30,
                    validation_split=0.2)
```

```
2022-03-20 12:29:48.141054: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 1242955776
exceeds 10% of free system memory.
2022-03-20 12:29:49.757808: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 1242955776
exceeds 10% of free system memory.
2022-03-20 12:29:50.753984: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR
Optimization Passes are enabled (registered 2)
```

Epoch 1/30

```
2022-03-20 12:29:52.907433: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369]
Loaded cuDNN version 8005
```

```
593/593 [=====] - 18s 17ms/step - loss: 16.1346 -
gender_out_loss: 0.6821 - age_out_loss: 15.4525 - gender_out_accuracy: 0.5472 -
age_out_accuracy: 0.0476 - val_loss: 12.7578 - val_gender_out_loss: 0.5521 -
val_age_out_loss: 12.2057 - val_gender_out_accuracy: 0.7269 -
val_age_out_accuracy: 0.0460
```

Epoch 2/30

```
593/593 [=====] - 9s 16ms/step - loss: 11.2216 -
gender_out_loss: 0.4761 - age_out_loss: 10.7455 - gender_out_accuracy: 0.7741 -
age_out_accuracy: 0.0285 - val_loss: 11.5279 - val_gender_out_loss: 0.4163 -
val_age_out_loss: 11.1116 - val_gender_out_accuracy: 0.8064 -
val_age_out_accuracy: 0.0255
```

Epoch 3/30

```
593/593 [=====] - 10s 16ms/step - loss: 9.3465 -
gender_out_loss: 0.3925 - age_out_loss: 8.9540 - gender_out_accuracy: 0.8214 -
age_out_accuracy: 0.0157 - val_loss: 8.4260 - val_gender_out_loss: 0.3558 -
val_age_out_loss: 8.0702 - val_gender_out_accuracy: 0.8361 -
val_age_out_accuracy: 0.0074
```

Epoch 4/30

```
593/593 [=====] - 9s 16ms/step - loss: 8.5809 -
gender_out_loss: 0.3446 - age_out_loss: 8.2363 - gender_out_accuracy: 0.8416 -
age_out_accuracy: 0.0119 - val_loss: 8.5080 - val_gender_out_loss: 0.3214 -
val_age_out_loss: 8.1866 - val_gender_out_accuracy: 0.8541 -
val_age_out_accuracy: 0.0078
```

Epoch 5/30

```
593/593 [=====] - 9s 16ms/step - loss: 8.0615 -
gender_out_loss: 0.3149 - age_out_loss: 7.7466 - gender_out_accuracy: 0.8602 -
age_out_accuracy: 0.0109 - val_loss: 7.5080 - val_gender_out_loss: 0.3134 -
```


val_age_out_loss: 7.1946 - val_gender_out_accuracy: 0.8551 -
 val_age_out_accuracy: 0.0076
 Epoch 6/30
 593/593 [=====] - 10s 16ms/step - loss: 7.6047 -
 gender_out_loss: 0.2935 - age_out_loss: 7.3112 - gender_out_accuracy: 0.8672 -
 age_out_accuracy: 0.0096 - val_loss: 7.5676 - val_gender_out_loss: 0.2822 -
 val_age_out_loss: 7.2854 - val_gender_out_accuracy: 0.8747 -
 val_age_out_accuracy: 0.0063
 Epoch 7/30
 593/593 [=====] - 9s 15ms/step - loss: 7.2906 -
 gender_out_loss: 0.2782 - age_out_loss: 7.0124 - gender_out_accuracy: 0.8743 -
 age_out_accuracy: 0.0080 - val_loss: 7.1280 - val_gender_out_loss: 0.2800 -
 val_age_out_loss: 6.8480 - val_gender_out_accuracy: 0.8739 -
 val_age_out_accuracy: 0.0049
 Epoch 8/30
 593/593 [=====] - 9s 16ms/step - loss: 6.9194 -
 gender_out_loss: 0.2654 - age_out_loss: 6.6540 - gender_out_accuracy: 0.8818 -
 age_out_accuracy: 0.0072 - val_loss: 8.0823 - val_gender_out_loss: 0.2770 -
 val_age_out_loss: 7.8053 - val_gender_out_accuracy: 0.8766 -
 val_age_out_accuracy: 0.0049
 Epoch 9/30
 593/593 [=====] - 9s 15ms/step - loss: 6.6902 -
 gender_out_loss: 0.2507 - age_out_loss: 6.4395 - gender_out_accuracy: 0.8903 -
 age_out_accuracy: 0.0064 - val_loss: 7.1591 - val_gender_out_loss: 0.2882 -
 val_age_out_loss: 6.8709 - val_gender_out_accuracy: 0.8707 -
 val_age_out_accuracy: 0.0032
 Epoch 10/30
 593/593 [=====] - 10s 16ms/step - loss: 6.4238 -
 gender_out_loss: 0.2404 - age_out_loss: 6.1834 - gender_out_accuracy: 0.8941 -
 age_out_accuracy: 0.0063 - val_loss: 7.0038 - val_gender_out_loss: 0.2649 -
 val_age_out_loss: 6.7389 - val_gender_out_accuracy: 0.8842 -
 val_age_out_accuracy: 0.0051
 Epoch 11/30
 593/593 [=====] - 10s 16ms/step - loss: 6.2591 -
 gender_out_loss: 0.2276 - age_out_loss: 6.0316 - gender_out_accuracy: 0.9011 -
 age_out_accuracy: 0.0063 - val_loss: 6.8535 - val_gender_out_loss: 0.2642 -
 val_age_out_loss: 6.5894 - val_gender_out_accuracy: 0.8876 -
 val_age_out_accuracy: 0.0027
 Epoch 12/30
 593/593 [=====] - 10s 16ms/step - loss: 5.9888 -
 gender_out_loss: 0.2179 - age_out_loss: 5.7709 - gender_out_accuracy: 0.9047 -
 age_out_accuracy: 0.0072 - val_loss: 6.8253 - val_gender_out_loss: 0.2690 -
 val_age_out_loss: 6.5562 - val_gender_out_accuracy: 0.8851 -
 val_age_out_accuracy: 0.0049
 Epoch 13/30
 593/593 [=====] - 10s 16ms/step - loss: 5.7775 -
 gender_out_loss: 0.2075 - age_out_loss: 5.5700 - gender_out_accuracy: 0.9118 -
 age_out_accuracy: 0.0059 - val_loss: 7.1583 - val_gender_out_loss: 0.2630 -



val_age_out_loss: 6.8953 - val_gender_out_accuracy: 0.8876 -
 val_age_out_accuracy: 0.0036
 Epoch 14/30
 593/593 [=====] - 9s 15ms/step - loss: 5.4795 -
 gender_out_loss: 0.1951 - age_out_loss: 5.2844 - gender_out_accuracy: 0.9160 -
 age_out_accuracy: 0.0054 - val_loss: 6.8055 - val_gender_out_loss: 0.2790 -
 val_age_out_loss: 6.5264 - val_gender_out_accuracy: 0.8838 -
 val_age_out_accuracy: 0.0034
 Epoch 15/30
 593/593 [=====] - 9s 16ms/step - loss: 5.4528 -
 gender_out_loss: 0.1831 - age_out_loss: 5.2697 - gender_out_accuracy: 0.9243 -
 age_out_accuracy: 0.0057 - val_loss: 6.9825 - val_gender_out_loss: 0.2813 -
 val_age_out_loss: 6.7012 - val_gender_out_accuracy: 0.8882 -
 val_age_out_accuracy: 0.0034
 Epoch 16/30
 593/593 [=====] - 9s 15ms/step - loss: 5.1696 -
 gender_out_loss: 0.1693 - age_out_loss: 5.0004 - gender_out_accuracy: 0.9304 -
 age_out_accuracy: 0.0051 - val_loss: 6.9286 - val_gender_out_loss: 0.2636 -
 val_age_out_loss: 6.6650 - val_gender_out_accuracy: 0.8817 -
 val_age_out_accuracy: 0.0044
 Epoch 17/30
 593/593 [=====] - 10s 16ms/step - loss: 5.0638 -
 gender_out_loss: 0.1638 - age_out_loss: 4.9000 - gender_out_accuracy: 0.9312 -
 age_out_accuracy: 0.0054 - val_loss: 6.9296 - val_gender_out_loss: 0.2813 -
 val_age_out_loss: 6.6483 - val_gender_out_accuracy: 0.8956 -
 val_age_out_accuracy: 0.0042
 Epoch 18/30
 593/593 [=====] - 9s 15ms/step - loss: 4.8813 -
 gender_out_loss: 0.1494 - age_out_loss: 4.7319 - gender_out_accuracy: 0.9396 -
 age_out_accuracy: 0.0053 - val_loss: 6.9294 - val_gender_out_loss: 0.2971 -
 val_age_out_loss: 6.6323 - val_gender_out_accuracy: 0.8880 -
 val_age_out_accuracy: 0.0040
 Epoch 19/30
 593/593 [=====] - 9s 15ms/step - loss: 4.8204 -
 gender_out_loss: 0.1428 - age_out_loss: 4.6776 - gender_out_accuracy: 0.9414 -
 age_out_accuracy: 0.0057 - val_loss: 6.9242 - val_gender_out_loss: 0.3056 -
 val_age_out_loss: 6.6185 - val_gender_out_accuracy: 0.8832 -
 val_age_out_accuracy: 0.0042
 Epoch 20/30
 593/593 [=====] - 9s 16ms/step - loss: 4.6624 -
 gender_out_loss: 0.1350 - age_out_loss: 4.5274 - gender_out_accuracy: 0.9444 -
 age_out_accuracy: 0.0057 - val_loss: 7.0920 - val_gender_out_loss: 0.3745 -
 val_age_out_loss: 6.7175 - val_gender_out_accuracy: 0.8699 -
 val_age_out_accuracy: 0.0070
 Epoch 21/30
 593/593 [=====] - 9s 16ms/step - loss: 4.5481 -
 gender_out_loss: 0.1267 - age_out_loss: 4.4214 - gender_out_accuracy: 0.9491 -
 age_out_accuracy: 0.0068 - val_loss: 6.9295 - val_gender_out_loss: 0.3286 -

val_age_out_loss: 6.6009 - val_gender_out_accuracy: 0.8865 -
 val_age_out_accuracy: 0.0032
 Epoch 22/30
 593/593 [=====] - 9s 16ms/step - loss: 4.4753 -
 gender_out_loss: 0.1239 - age_out_loss: 4.3514 - gender_out_accuracy: 0.9497 -
 age_out_accuracy: 0.0054 - val_loss: 7.0483 - val_gender_out_loss: 0.3409 -
 val_age_out_loss: 6.7075 - val_gender_out_accuracy: 0.8918 -
 val_age_out_accuracy: 0.0070
 Epoch 23/30
 593/593 [=====] - 9s 16ms/step - loss: 4.4120 -
 gender_out_loss: 0.1102 - age_out_loss: 4.3018 - gender_out_accuracy: 0.9540 -
 age_out_accuracy: 0.0090 - val_loss: 6.9948 - val_gender_out_loss: 0.3285 -
 val_age_out_loss: 6.6663 - val_gender_out_accuracy: 0.8895 -
 val_age_out_accuracy: 0.0105
 Epoch 24/30
 593/593 [=====] - 10s 16ms/step - loss: 4.2673 -
 gender_out_loss: 0.1059 - age_out_loss: 4.1614 - gender_out_accuracy: 0.9583 -
 age_out_accuracy: 0.0193 - val_loss: 7.0131 - val_gender_out_loss: 0.3328 -
 val_age_out_loss: 6.6803 - val_gender_out_accuracy: 0.8897 -
 val_age_out_accuracy: 0.0243
 Epoch 25/30
 593/593 [=====] - 9s 15ms/step - loss: 4.1578 -
 gender_out_loss: 0.1024 - age_out_loss: 4.0553 - gender_out_accuracy: 0.9582 -
 age_out_accuracy: 0.0264 - val_loss: 6.8706 - val_gender_out_loss: 0.3361 -
 val_age_out_loss: 6.5345 - val_gender_out_accuracy: 0.8958 -
 val_age_out_accuracy: 0.0287
 Epoch 26/30
 593/593 [=====] - 9s 15ms/step - loss: 4.0662 -
 gender_out_loss: 0.0933 - age_out_loss: 3.9730 - gender_out_accuracy: 0.9611 -
 age_out_accuracy: 0.0299 - val_loss: 7.2064 - val_gender_out_loss: 0.3738 -
 val_age_out_loss: 6.8326 - val_gender_out_accuracy: 0.8918 -
 val_age_out_accuracy: 0.0266
 Epoch 27/30
 593/593 [=====] - 9s 15ms/step - loss: 4.0040 -
 gender_out_loss: 0.0851 - age_out_loss: 3.9189 - gender_out_accuracy: 0.9644 -
 age_out_accuracy: 0.0311 - val_loss: 7.1397 - val_gender_out_loss: 0.4333 -
 val_age_out_loss: 6.7064 - val_gender_out_accuracy: 0.8903 -
 val_age_out_accuracy: 0.0331
 Epoch 28/30
 593/593 [=====] - 10s 16ms/step - loss: 3.9340 -
 gender_out_loss: 0.0848 - age_out_loss: 3.8492 - gender_out_accuracy: 0.9641 -
 age_out_accuracy: 0.0344 - val_loss: 7.0291 - val_gender_out_loss: 0.4004 -
 val_age_out_loss: 6.6287 - val_gender_out_accuracy: 0.8889 -
 val_age_out_accuracy: 0.0251
 Epoch 29/30
 593/593 [=====] - 9s 16ms/step - loss: 3.9378 -
 gender_out_loss: 0.0819 - age_out_loss: 3.8559 - gender_out_accuracy: 0.9659 -
 age_out_accuracy: 0.0329 - val_loss: 6.9958 - val_gender_out_loss: 0.3569 -

```
val_age_out_loss: 6.6389 - val_gender_out_accuracy: 0.8895 -  
val_age_out_accuracy: 0.0346  
Epoch 30/30  
593/593 [=====] - 9s 15ms/step - loss: 3.8026 -  
gender_out_loss: 0.0774 - age_out_loss: 3.7252 - gender_out_accuracy: 0.9671 -  
age_out_accuracy: 0.0341 - val_loss: 7.0322 - val_gender_out_loss: 0.4161 -  
val_age_out_loss: 6.6161 - val_gender_out_accuracy: 0.8920 -  
val_age_out_accuracy: 0.0259
```

0.6 Plot the Results

```
[27]: # plot results for gender  
acc = history.history['gender_out_accuracy']  
val_acc = history.history['val_gender_out_accuracy']  
epochs = range(len(acc))  
  
plt.plot(epochs, acc, 'b', label='Training Accuracy')  
plt.plot(epochs, val_acc, 'r', label='Validation Accuracy')  
plt.title('Accuracy Graph')  
plt.legend()  
plt.figure()  
  
loss = history.history['gender_out_loss']  
val_loss = history.history['val_gender_out_loss']  
  
plt.plot(epochs, loss, 'b', label='Training Loss')  
plt.plot(epochs, val_loss, 'r', label='Validation Loss')  
plt.title('Loss Graph')  
plt.legend()  
plt.show()
```




```
[28]: # plot results for age  
loss = history.history['age_out_loss']
```

```

val_loss = history.history['val_age_out_loss']
epochs = range(len(loss))

plt.plot(epochs, loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'r', label='Validation Loss')
plt.title('Loss Graph')
plt.legend()
plt.show()

```



1 Prediction with Test Data

```

[32]: image_index = 100
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])
# predict from model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');

```

Original Gender: Female Original Age: 3
 Predicted Gender: Female Predicted Age: 1



```
[33]: image_index = 3000
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])
# predict from model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Male Original Age: 28
Predicted Gender: Male Predicted Age: 32



```
[34]: image_index = 10000
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])
# predict from model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Male Original Age: 42
Predicted Gender: Male Predicted Age: 38

